

Branch: master ▾

r-cheat-sheet / Data Frames.md

Find file

Copy path



mattm Update notes

aaca237 on 2 May 2017

1 contributor

322 lines (243 sloc) 5.92 KB

Data Frames

A typical data set contains data of different modes. In an employee data set, for example, we might have character string data, such as employee names, and numeric data, such as salaries. So, although a data set of (say) 50 employees with 4 variables per worker has the look and feel of a 50-by-4 matrix, it does not qualify as such in R, because it mixes types.

Instead of a matrix, we use a data frame. A data frame in R is a list, with each component of the list being a vector corresponding to a column in our "matrix" of data. Indeed, you can create data frames in just this way:

The Art of R Programming

Creating a data frame from lists

```
> d <- data.frame(list(kids = c("Jack", "Jill"), ages = c(12, 10)))

> d
  kids ages
1 Jack   12
2 Jill   10
```

Checking the dimensions of a data frame

That's rows then columns:

```
> d <- data.frame(list(kids = c("Jack", "Jill", "Johnny"), ages = c(12, 10, 4)))
> d
  kids ages
1 Jack   12
2 Jill   10
3 Johnny  4
> dim(d)
[1] 3 2
```

Returning the column names

```
> colnames(d)
[1] "kids" "ages"
```

Viewing a summary of the data

```
> summary(d)
      kids      ages
Jack  :1   Min.   : 4.000
Jill  :1   1st Qu.: 7.000
Johnny:1   Median :10.000
      Mean    : 8.667
      3rd Qu.:11.000
      Max.    :12.000
```

Viewing the structure of the data

```
> str(d)
'data.frame':   3 obs. of  2 variables:
 $ kids: Factor w/ 3 levels "Jack","Jill",...: 1 2 3
 $ ages: num  12 10 4
```

Returning the values of a data frame component

Just like you would a list:

```
> d$kids
[1] Jack Jill
Levels: Jack Jill
```

Returning a component of the data frame

Use single brackets [] to return a list:

```
> d['kids']
      kids
1 Jack
2 Jill
```

Using the standard [] method

```
> d <- data.frame(list(kids = c("Jack", "Jill"), ages = c(12, 10)))
> d[d$kids == "Jack",]
      kids ages
1 Jack    12
```

Subsetting using subset

Single column, exact value

```
> housing <- read.csv("data/landdata-states.csv")
```

With subset :

```
> fl = subset(housing, State == "FL")
```

With dplyr 's filter function:

```
> fl = filter(housing, State == "FL")
```

Single column, any of multiple values

With subset :

```
both = subset(housing, State %in% c("FL", "GA"))
```

With dplyr's filter :

```
> both = filter(housing, State == "FL" | State == "GA")
```

Multiple columns

With subset :

```
> subset(housing, State == "AK" & Home.Value == 224952)
```

With dplyr's filter :

```
> filter(housing, State == "AK" & Home.Value == 224952)
```

Re-ordering rows

With order :

```
> # Ascending
> housing[order(housing$Home.Value), ]

> # Descending
> housing[order(-housing$Home.Value), ]
> housing[order(housing$Home.Value, decreasing = TRUE), ]
```

With dplyr's arrange :

```
> # Ascending
> arrange(housing, Home.Value)

> # Descending
> arrange(housing, desc(Home.Value))
```

Selecting columns

With subsetting:

```
> housing[, c("State", "Home.Value")]
```

With dplyr's select :

```
> select(housing, State, Home.Value)
```

Removing a column

With `select` :

```
housing <- select(housing, -State)
```

Renaming a column

With dplyr's `rename` :

```
> rename(housing, State.Name = State)
```

Note that `State.Name` is the *new* name.

Extract distinct (unique) rows

With `unique` :

```
> unique(housing[, c("State", "region")])
```

With dplyr's `distinct` :

```
> distinct(housing, State, region)
```

Removing NA values

Use `complete.cases` :

```
> d <- data.frame(list(kids = c("Jack", "Jill"), ages = c(12, NA)))
> d
  kids ages
1 Jack  12
2 Jill  NA
> d[complete.cases(d), ]
  kids ages
1 Jack  12
```

Taking a sample

With `sample` :

```
> housing[sample(nrow(housing), 5), ]
```

([Source](#))

With dplyr's `sample_n` :

```
> sample_n(housing, 5)
```

There's also a `sample_frac` that grabs a random percentage.

Adding a new column

```
> d <- data.frame(list(name = c("Jack", "Jill"), age = c(12, 10)))
> d
  name age
1 Jack  12
2 Jill  10
```

Natively:

```
> d$next_age = d$ages + 1
> d
  name age next_age
1 Jack  12       13
2 Jill  10       11
```

With dplyr's `mutate` :

```
> d <- mutate(d, next_age = ages + 1)
> d
  name age next_age
1 Jack  12       13
2 Jill  10       11
```

Note that with `mutate` , you can reference columns you're currently adding:

```
> d <- mutate(d, next_age = age + 1, next_next_age = next_age + 1)
> d
  name age next_age next_next_age
1 Jack  12       13           14
2 Jill  10       11           12
```

Grouping Operations

Applying `summarize` to groups of observations

With one summary statistic:

```
> by_state = group_by(housing, State)
> summarize(by_state, Avg.Home.Value = mean(Home.Value))
# A tibble: 51 × 2
  State      Avg.Home.Value
  <fctr>      <dbl>
1 AK      147385.14
2 AL      92545.22
3 AR      82076.84
4 AZ      140755.59
```

With multiple:

```
> by_state = group_by(housing, State)
> summarize(by_state, count = n(), Avg.Home.Value = mean(Home.Value))
# A tibble: 51 × 3
  State count Avg.Home.Value
  <fctr> <int>      <dbl>
1 AK    153      147385.14
2 AL    153      92545.22
3 AR    153      82076.84
```

4	AZ	153	140755.59
5	CA	153	282808.08

Note that `n()` is an *aggregate function* provided by `dplyr` that returns the number of observations in each group, which in this example are all 153.

There are other aggregate functions as well: `n_distinct(x)` (the number of unique values of `x`), `first(x)`, `last(x)`, and `nth(x)`.

We can chain `dplyr` functions together using the `%>%` operator:

```
> group_by(housing, State) %>% summarize(Avg.Home.Value = mean(Home.Value))
# A tibble: 51 × 2
  State Avg.Home.Value
  <fctr>      <dbl>
1     AK      147385.14
2     AL      92545.22
3     AR       82076.84
4     AZ      140755.59
5     CA      282808.08
```