

Tree based methods LAB

Enrique J. De La Hoz D.

Data Science UTB

Building a simple decision tree

The *clientes* dataset contains information about the record of clients in a telecommunication company.

You will use a decision tree to try to learn patterns in the **churn rate**, based on the tenure and the Monthly Payment.

Then, see how the tree's predictions differ for a client leaving the company versus one remaining as an active client.

- Load 'clients.csv'
- Create a training and test dataset.
- Load the rpart package.
- Fit a decision tree model using the training dataset with the function `rpart()`, call it `clients_model`
 - Supply the R formula that specifies outcome as a function of *tenure* and *Pago_mensual* as the first argument.
 - Leave the control argument alone for now. (You'll learn more about that later!)
- Use `predict()` with the resulting loan model to predict the outcome for the churn rate. Use the type argument to predict the "class" of the outcome.
- Apply the model to the test dataset and check the results

Visualizing classification trees

- As Data Scientist it is important to support the results through visualizations
- The structure of classification trees can be depicted visually, which helps to understand how the tree makes its decisions.
- Type `clients_model` to see a text representation of the classification tree.
- Load the `rpart.plot` package.
- Apply the `rpart.plot()` function to the `clients_model` to visualize the tree.
- See how changing other plotting parameters impacts the visualization by running the supplied command.

Building and evaluating a larger tree

Previously, you created a simple decision tree that used the client's information to predict the Churn.

Using all of the available applicant data, build a more sophisticated lending model

- The rpart package is loaded into the workspace and the `loans_train` and `loans_test` datasets have been created.
- Use `rpart()` to build a loan model using the training dataset and all of the available predictors. Again, leave the control argument alone.
- Applying the `predict()` function to the testing dataset, create a vector of predicted outcomes. Don't forget the type argument.

- Create a table() to compare the predicted values to the actual outcome values.
- Compute the accuracy of the predictions using the mean() function.

Preventing overgrown trees

The tree grown on the full set of clients data grew to be extremely large and extremely complex, with hundreds of splits and leaf nodes containing only a handful of applicants. This tree would be almost impossible for a loan officer to interpret (Difficult to sell).

Using the pre-pruning methods for early stopping, you can prevent a tree from growing too large and complex. See how the rpart control options for maximum tree depth and minimum split count impact the resulting tree.

- Add a maxdepth parameter to the rpart.control() object to set the maximum tree depth to six. Leave the parameter cp = 0. Pass the result of rpart.control() as the control parameter in your rpart() call.
- Check how the test set accuracy of the simpler model compares to the original accuracy.
- First create a vector of predictions using the predict() function.
- Compare the predictions to the actual outcomes and use mean() to calculate the accuracy.
- Add a minsplit parameter to the rpart.control() object to require 70 observations to split. Again, leave cp = 0.
- Again compare the accuracy of the simpler tree to the original.

Creating a nicely pruned tree

Stopping a tree from growing all the way can lead it to ignore some aspects of the data or miss important trends it may have discovered later.

By using post-pruning, you can intentionally grow a large and complex tree then prune it to be smaller and more efficient later on.

In this exercise, you will have the opportunity to construct a visualization of the tree's performance versus complexity, and use this information to prune the tree to an appropriate level.

The rpart package is loaded into the workspace, along with loans_test and loans_train.

- Use all of the applicant variables and no pre-pruning to create an overly complex tree. Make sure to set cp = 0 in rpart.control() to prevent pre-pruning.
- Create a complexity plot by using plotcp() on the model.
- Based on the complexity plot, prune the tree to a complexity of 0.0014 using the prune() function with the tree and the complexity parameter.
- Compare the accuracy of the pruned tree to the original accuracy of 58.3%. - To calculate the accuracy use the predict() and mean() functions